


Fig.\,1. AST INT8 deployment pipeline. A single quantized checkpoint admits AST-class capability into edge tiers (left, green) previously confined to CNN-scale models, while simultaneously enabling 4× multi-tenant density on cloud GPUs (right).



Same INT8 checkpoint runs in all tiers; each device picks the backend (CoreML / ONNX-RT ARM / TFLite / TensorRT / fbgemm).